

Introduction to Statistical Machine Learning



Xiang Zhou

School of Data Science
Department of Mathematics
City University of Hong Kong

Notations

- input - output relation
 - ▶ x : inputs, feature vectors, predictors, independent variables.
 - ▶ y : output, response, dependent variable.
- \mathcal{X} and \mathcal{Y} denote the spaces of the generic x and y variables, respectively.
 - ▶ Generally $\mathcal{X} = \mathbb{R}^p$; qualitative features are coded using, for example, dummy variables (such as 0,1, -1, etc).
 - ▶ Typically $\mathcal{Y} \in \mathbb{R}^1$ is a scalar, or takes a finite number of values as a subset of \mathbb{N} ; it can be a vector in some scenarios.
- The random variable (X, Y) has the joint distribution $p(x, y)$ on the sample space $\mathcal{X} \times \mathcal{Y}$.

- It is usually assumed that the ground truth for the relation between from input to the output is a deterministic input-output mapping from $x \in \mathcal{X}$ to $y_{\text{true}} \in \mathcal{Y}$:

$$y_{\text{true}} = f^*(x)$$

where the ground truth f^* is an unknown function and has to be approximated by learning from the training dataset.

- In **supervised learning**, the data (observations) are given as the collection of the pairs ¹

$$\mathcal{D} = \left\{ (x^{(i)}, y^{(i)}) : 1 \leq i \leq N \right\} \subset (\mathcal{X} \times \mathcal{Y})^N$$

which is assumed iid samples of the r.v. (X, Y) with an unknown joint distribution $p(x, y)$ on the product space $\mathcal{X} \times \mathcal{Y}$.

- ▶ **Regression**: \mathcal{Y} is continuous/quantitative, e.g., \mathbb{R}^d or its subset.
- ▶ **Classification**: \mathcal{Y} is discrete and finite (categorical variable), encoded by a finite number $\{1, \dots, K\}$. In this case, “ y ” is usually called “label”.
- In **unsupervised learning**, the observations only have $\{x^{(i)}\}$, the information $y^{(i)}$ is zero or there is no definition of y variable. The purpose is to identify the pattern of $\{x^{(i)}\}$ itself, such as model reduction.

¹sometimes it is denoted as $\mathcal{D} = \{(x_i, y_i) : 1 \leq i \leq n\}$ if the subindex has no ambiguity.

Output as perturbed truth

- The observed $x^{(i)} \in \mathcal{X}$ are samples from the marginal distribution p_X , i.e., $x^{(i)} \sim X$; in some cases, they are deterministic and assigned by a procedure of experiment design.
- The corresponding observed $y^{(i)}$ are assumed to be the *perturbed* truth $f^*(x^{(i)})$ with additional measurement error $\varepsilon^{(i)}$ which are assumed to be iid and independent from X .

$$y^{(i)} = f^*(x^{(i)}) + \varepsilon^{(i)}.$$

$\{\varepsilon^{(i)}\}$ are assumed iid and distributed as a generic r.v. ε .

- The effect of the noise ε can never be eliminated by any statistical learning algorithms (**irreducible error**).

- So, the joint distribution $p(x, y)$ of (X, Y) is completely determined by the triplet:

$$(p_X, f^*, p_\varepsilon)$$

- ▶ p_X : the distribution of the input
 - ▶ f^* : the input-output function,
 - ▶ p_ε : the distribution of error.
- The joint distribution $p(x, y)$ manifests by the available dataset D . Given X, Y random variables $\sim p$, how to identify f^* ?

Statistics and machine learning

Different views and terminologies:

Machine Learning	Statistics
Supervised learning	Classification/regression
Unsupervised learning	Clustering
Semisupervised learning	Classification/regression with missing responses
Features/outcomes	Covariates/responses
Training set/testing set	Sample/population
Learner	Statistical model
Generalization error	Misclassification error/prediction error

Suggestion

Learn machine learning like a statistician or an applied mathematician, not a software engineer

- Start from stating a problem, not show an algorithm first: Many times, the students think the methods/algorithms/procedures as the problem itself.
- Pay close attentions to the “modelling” process: how to turn the data problem into the statistical model. In particularly the underlying principle which applies very general.
- Try to rigorously (and in most general context) understand (like a math theorem) the heuristic arguments used in practice, even for some toy examples.
- Draw strict boundaries between general principles and specific methods. In between, computational method survives and thrives.
- Diagnose and rationalize your results of numerical experiments. Test different dataset/parameters/methods.

Learning Theory: An Approximation Theory Viewpoint

Reference:

Learning Theory: An Approximation Theory Viewpoint; by F. Cucker and D.X. Zhou, Cambridge University Press 2007

Given r.v.s X and Y , find a function $f : \mathcal{X} \rightarrow \mathcal{Y}$ so that $f(X)$ can explain Y best in certain sense.

Conditional Expectation as Optimal Prediction

The best L^2 approximation of a function f of the r.v. X to a r.v. Y is achieved by the conditional probability. The **(generalized) squared error**¹

$$\mathcal{E}(f) := \mathbb{E} |Y - f(X)|^2 \quad (1)$$

has a minimum at

$$f^*(x) = \mathbb{E}(Y|X = x).$$

i.e.,

$$\mathbb{E} |Y - f^*(X)|^2 = \min_{f: \text{ a Borel function}} \mathbb{E} (|Y - f(X)|^2)$$

f in fact should satisfy the condition that the RHS L_2 norm exists.

¹also named as generalization risk, mean square error, L_2 error, etc.
Xiang Zhou CityU

Proof.

- We show first that $\mathbb{E}[(Y - f^*(X))h(X)] = 0$ ^a is true for any function h . Using the double expectation theorem^b, we have

$$\begin{aligned}\mathbb{E}[(Y - f^*(X))h(X)] &= \mathbb{E}[\mathbb{E}[Y - f^*(X)|X]h(X)] \\ &= \mathbb{E}[\mathbb{E}(Yh(X)|X) - f^*(X)h(X)] = 0.\end{aligned}$$

- Note that

$$(y - f(x))^2 = (y - f^*(x))^2 + (f(x) - f^*(x))^2 - 2(y - f^*(x))h(x)$$


where $h(x) = f(x) - f^*(x)$, then for any f

$$\boxed{\mathbb{E}(|f(X) - Y|^2) = \mathbb{E}(|f^*(X) - Y|^2) + \mathbb{E}[|f(X) - f^*(X)|^2]} \quad (2)$$



^asometimes it is denoted $Y - f^*(X) \perp h(X)$, the perpendicular property in L_2 space.

^b $\mathbb{E}[\mathbb{E}(Y|X)] = \mathbb{E}Y$



reference for elementary math: [Understanding Conditional Expectation via Vector Projection](#)

The following exercise is to use the viewpoint of variational calculus: minimizing functions in the function space

Exercise

Use the method of Calculus of Variation to solve ^a

$$\inf_f \iint (f(x) - y)^2 p_{X,Y}(x, y) dx dy$$

where $p_{X,Y}$ is the joint pdf of the r.v.s (X, Y) . The optimal f^ satisfies*

$$\int (f^*(x) - y) p_{X,Y}(x, y) dy = 0, \quad \forall x$$

i.e., $\mathbb{E} Y = \mathbb{E} f^(X)$*

^aRigorously, f is in the p -weighted L_2 space

What if generalized the L_2 norm to L_p norm ?

Exercise (Conditional distribution of multivariate Gaussian r.v.)

Suppose that $X = (X_1, X_2)$ is a two dimensional Gaussian random variable with mean $\mu = (\mu_1, \mu_2)$ and the covariance matrix

$\Sigma = \begin{pmatrix} \sigma_1^2 & \rho\sigma_1\sigma_2 \\ \rho\sigma_1\sigma_2 & \sigma_2^2 \end{pmatrix}$. What is the conditional pdf $p(x_1|x_2)$ of X_1 given $X_2 = x_2$? What is the expectation of X_1 given $X_2 = x_2$? What is the **general result** in n dimension ?

$$X_1|_{X_2=x_2} \sim \mathcal{N}(\mu_1 + \rho\frac{\sigma_1}{\sigma_2}(x_2 - \mu_2), \sigma_1^2(1 - \rho^2))$$

- The objective function $\mathcal{E}(f)$ is also called loss function, risk, etc.
- $\mathbb{E} f^*(X) = \mathbb{E} Y$: $f^*(X)$ is an unbiased estimate of Y ;
- The variance of the difference between Y and the optimal prediction $f^*(X)$ is

$$\sigma_*^2(x) := \mathbb{E} [(Y - f^*(X))^2 | X = x]$$

- Take average of $\sigma^2(x)$ over x , then the averaged uncertainty is

$$\sigma_*^2 := \mathbb{E}_X \sigma^2(X) = \mathbb{E} [|Y - f^*(X)|^2] = \mathcal{E}(f^*)$$

is the variance of the measurement error $Y - f^*(X)$: **irreducible error**

We have shown that (2) for any two r.v. X, Y :

$$\mathcal{E}(f) = \underbrace{\mathbb{E}(|f(X) - Y|^2)}_{\text{Mean Square Error}} = \underbrace{\mathbb{E}_{X,Y}(|f^*(X) - Y|^2)}_{\mathcal{E}(f^*), \text{Bayes error}} + \underbrace{\mathbb{E}_X [|f(X) - f^*(X)|^2]}_{\text{model error}} \quad (3)$$

where $f^*(x) = \mathbb{E}(Y|X = x)$ is called **Bayes rule**.

- Bayes error: irreducible error;
- Model error: the distance from f to f^* .

Application to classification with 0-1 loss

- Assume $Y \in \{1, \dots, K\}$ and $X \in \mathbb{R}^p$.
- $f : \mathcal{X} \rightarrow \mathcal{Y}$ is a piece-wise constant function.
- A Loss function $\ell(Y, f(Y))$ for penalizing errors in misclassification.
- Most Common choice is the 0-1 loss

$$\ell(Y, f(X)) = I(Y \neq f(X)) = \begin{cases} 1 & \text{if } Y \neq f(X) \\ 0 & \text{if } Y = f(X) \end{cases}.$$

- The expected prediction error, or the generalization error, is

$$\mathcal{E}(f) = \mathbb{E} I(Y \neq f(X)) = \mathbb{P}(Y \neq f(X)) = 1 - \mathbb{P}(Y = f(X)).$$

- The Bayes rule minimizing $\mathcal{E}(f)$ is the one maximizing $\mathbb{P}(Y = f(X)) = \int_{\mathcal{X}} \mathbb{P}(Y = f(x)|X = x)p_X(x)dx$, which is

$$f^*(x) = \underset{k}{\operatorname{argmax}} \mathbb{P}(Y = k|X = x).$$

since $\mathbb{P}(Y = k|X = x) \leq \mathbb{P}(Y = f^*(x)|X = x), \forall k$

How do we estimate f^* ?

- direction method by definition of conditional expectation;
- optimization method by minimization of generalization error;

Nearest neighbors method to approximate condition expectation

A natural way to approximate the function $\mathbb{E}(Y|X = x)$ is to replace the expectation by the average:

$$f^*(x) = \mathbb{E}(Y|X = x) \approx \text{AVE} \{y_i : \text{where } x_i = x\}, \quad \forall x \in \mathcal{X}$$

or

$$f^*(x) = \int_{\mathcal{Y}} y p_{Y|X}(y; x) dy$$

where $p_{Y|X}(y; x) = \frac{p_{X,Y}(x,y)}{p_X(x)}$ by estimating the density $p_{X,Y}(x, y)$ first?

Question: Would these work well in practice ? in high dimension ?

(1) this formula is defined point-wisely; (2) one does not have an accurate estimation of the expectation or the joint distribution $p(x, y)$, particularly in high dimension.

- Use windows with size Δ

$$\mathbb{E}(Y|X = x) \approx \text{AVE} \{y_i : \text{where } |x_i - x| \leq \Delta\}$$

- Use k -nearest neighbors: have a look at its neighbors, and take a vote:

$$\mathbb{E}(Y|X = x) \approx \text{AVE} \{y_i : x_i \in \mathcal{N}_k(x)\}$$

where $\mathcal{N}_k(x)$ is a neighborhood of x that contains exactly k neighbors (k -nearest neighbors).

Curse of dimensionality

k -NN can fail in high dimensions, because it becomes difficult to gather k observations close to a target point x_0 .

- Neighborhoods tend to be spatially large, and estimates are biased.
- Reducing the spatial size of the neighborhood means reducing k , and the variance of the estimate increases.

In general (see Figure 2.6),

- Most points are at the boundary, and points tend to be equidistant (Hall et al., JRSS-B, 2005).
- Sampling density is proportional to $n^{1/p}$: the number of sample size increases exponentially in dimension p : If 100 points are sufficient to estimate a function in \mathbb{R}^1 , 100^{10} are needed to achieve similar accuracy in \mathbb{R}^{10} .

This is similar to the curse of dimensionality in using the mesh grid based method to solve high dimensional PDE.

Minimizing the generalization error: learning as minimization

- Loss function $\ell(y, \hat{y}) : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}^+ \cup \{0\}$
 - ▶ L_2, L_1 norm
 - ▶ 0-1 loss: $\ell(y, \hat{y}) = I(y \neq \hat{y})$.
 - ▶ ...
- Population loss(Generalization error): $\mathcal{E}(f) = \mathbb{E} \ell(Y, f(X))$

The direct approach of minimizing the risk

$$\min_f \mathcal{E}(f) = \mathbb{E} \ell(Y, f(X))$$

where f can be in a very general class of functions, such as continuous function on \mathcal{X} , or even piecewise continuous function.

It is important to specify a hypothesis space \mathcal{H} to restrict the search space.

- ① the function space of f to search: hypothesis space(model class) \mathcal{H} .
- ② approximate the expectation \mathbb{E} in \mathcal{E} with the use of data $D = \{x_i, y_i\}$.

- ① The minimizer of

$$\min_{f \in \mathcal{H}} \mathcal{E}(f)$$

is denoted by $f_{\mathcal{H}}$.

- ② Population risk is approximated by empirical risk associated with $D = \{(x_i, y_i)\} : i = 1, \dots, N$:

$$\mathcal{E}(f) = \mathbb{E} \left[|Y - f(X)|^2 \right] \approx \mathcal{E}_N(f) = \frac{1}{N} \sum_i \ell(y_i, f(x_i))$$

The learning algorithm produces the learnt function \hat{f}_D (**regression function**) as the solution to

$$\mathcal{E}_N(\hat{f}_D) = \min_{f \in \mathcal{H}} \mathcal{E}_N(f) = \min_{f \in \mathcal{H}} \frac{1}{N} \sum_{(x_i, y_i) \in D} \ell(y_i, f(x_i))$$

- ▶ The **training error** is $\mathcal{E}_N(\hat{f}_D)$.
- ▶ The **test error** is $\mathcal{E}(\hat{f}_D)$.

- \hat{f}_D depends on *three* elements: the hypothesis space \mathcal{H} , the (training) data D (random !) and the loss function ℓ . So, \hat{f}_D is a random function to approximate the ground truth f .
- For the fixed \mathcal{H} and ℓ , \hat{f}_D is essentially a mapping from D to the function space \mathcal{H} . As $N \rightarrow \infty$, $\hat{f}_D \rightarrow f_{\mathcal{H}}$ by the law of large number.

Hypothesis space: representation of the function f

- parametric approach: e.g. $f(x) = \theta \cdot x$ with a set of finite and (usually fixed) number of parameters θ
- non-parametric approach: functional analysis viewpoint, f is only in some function space. (such as in the traditional computational math for representing the solution to some PDE)
- there is no rigorous boundary between parametric/non-parametric: eventually the computer represents a function f only by a finite number of freedoms.
- need to be restricted to certain model class in practice. Examples: linear model, polynomial, spline, kernel machine, neural network, etc.

some examples of \mathcal{H}

- Linear function (chapter 3 , ESL)
- Basis functions and dictionary methods (chapter 5, ESL)

$$f_{\theta}(x) = \sum_{m=1}^M \theta_m h_m(x),$$

where h_m 's are pre-specified basis functions.

- Kernel methods and local regression (chapter 6, ESL)

$$RSS(f, x_0) = \sum_{i=1}^n K(x_0, x_i)(y_i - f(x_i))^2,$$

where K is a kernel measuring the closeness between points.

- Roughness penalty (chapter 5, ESL)

$$\text{Penalized RSS}(f; \lambda) = \text{RSS}(f) + \lambda J(f),$$

where $J(f)$ is a regularization term on the model complexity.

The generalization error is then decomposed into

$$\mathcal{E}(f^*) = \underbrace{\mathcal{E}(f^*) - \mathcal{E}(f_{\mathcal{H}})}_{\text{approximation error}} + \underbrace{\mathcal{E}(f_{\mathcal{H}}) - \mathcal{E}_N(\hat{f}_{\mathcal{D}})}_{\text{sampling error}} + \underbrace{\mathcal{E}_N(\hat{f}_{\mathcal{D}})}_{\text{training error}}$$

- Approximation error: characterized how large (complex) the hypothesis space \mathcal{H} is. analogy to the interpolation inequality in finite element method. A larger space of \mathcal{H} indicates
 - ▶ A larger effective dimension of this space;
 - ▶ A smaller approximation error; lower bias
 - ▶ More difficult for numerical optimization;
- Sampling error: vanishes if the number of data points N goes to infinity. This error may be analyzed by probability inequality (e.g. Chebychev); analysis method used for Monte Carlo simulation. A larger sample size for training data indicates:
 - ▶ Smaller sampling error; lower variance;
 - ▶ more challenges for data collection and Big Data techniques;
- Training error:
 - ▶ design of optimization method, numerical convergence rate,
 - ▶ the time and memory cost ...
 - ▶ the output is a further approximation to \hat{f}_D : the optimization algorithm stops with a time cost T : $\hat{f}_D^{(T)}$

Very Important Remarks

- Although different fields focus on each of three errors separately, to address the problems for machine learning needs trans-disciplinary efforts:
 - ▶ approximation theory
 - ▶ statistical learning
 - ▶ computation

A joint strategy is very important to have a complete picture of the machine learning problem.

- These three topics are not completely separate. For example, the necessary number of sample size N may depend on the complexity of hypothesis space, which further depends on the dimension of \mathcal{X}, \mathcal{Y} and the specific representation form of functions in \mathcal{H} .

The optimal function in theory is f^* ; the optimal function in practical computation is \hat{f}_D . The difference in risks of these two:

$$\mathcal{E}(\hat{f}_D) - \mathcal{E}(f^*)$$

is called **generalization gap**: it measures to which extent the performance of the algorithm learnt from a given training data set D is valid for the whole distribution $p(x, y)$.

First inequality for Generalization Gap

Exercise

Prove the upper bound of the generalization gap:

$$\begin{aligned}\mathcal{E}(\hat{f}_D) - \mathcal{E}(f^*) &\leq 2 \sup_{f \in \mathcal{H}} |\mathcal{E}_N(f) - \mathcal{E}(f)| \\ &\approx 2 \sup_{f \in \mathcal{H}} \frac{1}{\sqrt{N}} \mathbb{V}(\ell(Y, f(X)))\end{aligned}$$

The first RHS depends on the hypothesis space \mathcal{H} and the dataset D of size N in \mathcal{E}_N ; and the second depends on \mathcal{H} and the joint distribution $p(x, y)$ of (X, Y) .

The supremum in the above inequality can describe the **complexity** of hypothesis space \mathcal{H} , which increases when \mathcal{H} is “larger”. WARNING: this upper bound is not sharp. A good lower bound is extremely difficult to obtain.

Statistical Learning: Viewpoint of Bias-Variance Trade-off

Let f be a given deterministic function $\mathcal{X} \rightarrow \mathcal{Y}$.

Assume that the response r.v. Y is defined by

$$Y = f(X) + \varepsilon$$

where the r.v. ε is independent of X and has mean 0 and variance σ_ε^2 .

Then $f^*(x) = \mathbb{E}[Y|X = x] = f(x)$ is the optimal approximation in the square error sense.

test error decomposition

Based on a training dataset $D = \{x_i, y_i\}$ where $y_i = f(x_i) + \varepsilon_i$ and $\varepsilon_i \sim \varepsilon$, one learns a regression function, denoted by $\hat{f}_D(\cdot)$, then the **test error**¹ at a new input x_0 (with a new independent measurement error $\varepsilon_0 \sim \varepsilon$) is

$$\begin{aligned}\mathbb{E}[(Y - \hat{f}_D(X))^2 | X = x_0] &= \mathbb{E}_{\varepsilon_0, D}[(f(x_0) + \varepsilon_0 - \hat{f}_D(x_0))^2] \\ &= \sigma_\varepsilon^2 + \mathbb{E}_D[(f(x_0) - \hat{f}_D(x_0))^2] \quad \because \mathbb{E}(\varepsilon_0) = 0, \text{ and } D \perp \varepsilon_0 \\ &= \sigma_\varepsilon^2 + \mathbb{E}_D[(f(x_0) - \mathbb{E} \hat{f}_D(x_0) + \mathbb{E} \hat{f}_D(x_0) - \hat{f}_D(x_0))^2] \\ &= \underbrace{\sigma_\varepsilon^2 + (f(x_0) - \mathbb{E} \hat{f}_D(x_0))^2}_{\text{Bias}^2} + \underbrace{\mathbb{V}(\hat{f}_D(x_0))}_{\text{Variance}}\end{aligned}$$

Sometimes, this quantity should be taken expectation w.r.t. $x_0 \sim p_X$.

¹also called **Prediction Error**. In some cases, this definition does not contain ε_0 contribution.

- low bias: large model space, low training error, overfitting, bad generalization ability (high variance);
- low variance: rigid model space, insensitive to the perturbation of the dataset used in fitting; good generalization for the new data from the same distribution.
- BAD news: it is almost impossible to decrease the bias and variance terms simultaneously!
- Criteria for model assessment or variable selection: **good trade-off between the bias and variance**
 - ▶ Analyse the bias or variance or model complexity to have analytical results (very limited cases)
 - ▶ A practical method is to estimate $\mathbb{E}[(Y - \hat{f}_D(X))^2] \approx \frac{1}{n} \sum_j (Y'_j - \hat{f}_D(X'_j))^2$ where X'_j are not from D but an independent dataset D' .
 - ▶ Where is the “new” dataset D' ? cross-validation (split 1 cent into 1/2 cents) and bootstrap (1 cent used twice) .




Figure: Bias-Variance Tradeoff as a Function of Model Capacity/Complexity

Typical objectives:

- Choose a value of a tuning parameter (**hyperparameter**) used in the model (such as k in k-NN)
- Estimate the prediction performance (test error) of a given model

Remarks:

- For both objectives, the best approach is to run the procedure on an independent test set, if one is available.
- If possible, one should use different test data for (1) and (2) above: a *validation set* for (1) and a *test set* for (2).
- Often there is insufficient data to create a separate validation or test set. In this case, *Cross-Validation* is useful.

K-fold cross validation

Denote the hyper-parameter by λ . K-fold cross validation is the most popular method for estimating a tuning parameter λ .

Divide the dataset (of size N) into K subsets: $\mathcal{A}_1, \dots, \mathcal{A}_K$ ($K = 2, 5, 10$ or N)

- For each $k = 1, \dots, K$, fit the model with parameter λ to $\{\mathcal{A}_1, \dots, \mathcal{A}_{k-1}, \mathcal{A}_{k+1}, \dots, \mathcal{A}_K\}$ giving $f_{\lambda}^{-k}(\cdot)$, and compute its prediction error on \mathcal{A}_k :

$$E_k(\lambda) = \sum_{x_i \in \mathcal{A}_k} \ell(y_i, f_{\lambda}^{-k}(x_i)).$$

- The average of these K values $E_k(\lambda)$ give the cross-validation error (per sample)

$$CV(\lambda) := \frac{1}{N} \sum_{k=1}^K E_k(\lambda).$$

- Choose the optimal λ^* yielding the smallest $CV(\lambda)$.

K-fold cross validation

- Cross-validation is often abbreviated as CV.
- In the subset selection procedure, λ is the subset size
- $f^{-k}(\lambda)$ is the best model of size λ , found from the training set that leaves out the k -th part of the data
- $E_k(\lambda)$ is its estimated test error on the k -th part.
- Using K -fold CV, the K test error estimates are averaged to give the final CV estimated test error.
- The output is the model associated with λ^* , typically, computed by using all N data.

- Bootstrap works by sampling N times with replacement from the training set to form a “bootstrap” data set. Then model is estimated on the bootstrap data set, and predictions are made on the original training set.
- This process is repeated many times and the results are averaged.
- *Bootstrap is most useful for estimating standard errors of predictions.*
- Can also use modified versions of the bootstrap to estimate prediction error. Sometimes produces better estimates than CV (still an open question!)